

Spark Developer Training - 3 Days

Manaranjan Pradhan

manaranjan@enablecloud.com

This notebook is given as part of Spark Training to Participants. Forwarding others is strictly prohibited.

Lab: Monitoring and Debugging Spark Applications

Things to learn

- Use of HDFS UI
- Use of Resource Manager UI
- Use of Application Monitoring UI
- Understand Execution Workflow

In [2]:

```
sc
```

Out[2]:

```
<pyspark.context.SparkContext at 0x7f639419a390>
```

Open <http://hadooplab.bigdataleap.com:50070/> (<http://hadooplab.bigdataleap.com:50070/>) on your browser for HDFS UI

This Spark Application is running on YARN. So, open the YARN Resource manager UI and verify if the application is running

- To open resource manager web UI, enter <http://hadooplab.bigdataleap.com:8088/> (<http://hadooplab.bigdataleap.com:8088/>)

Open the <http://hadooplab.bigdataleap.com:4040/> (<http://hadooplab.bigdataleap.com:4040/>) (or subsequent port) on the VM in firefox browser for application monitoring UI

Initialize SQLContext from SparkContext

In [3]:

```
from pyspark.sql import SQLContext
sqlContext = SQLContext(sc)
```

Reading JSON file from HDFS

In [4]:

```
## Read the json file from HDFS
txns = sqlContext.read.json( "hdfs://sparklab.awesomestats.in/sparklab/txnjsonsmall")
```

Go to the stages tab

- The data is read

Spark1.6.0

Jobs

Stages

Storage

Environment

Executors

SQL

Stages for All Jobs

Completed Stages: 1

Completed Stages (1)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Input
0	<div>json at NativeMethodAccessorImpl.java:-2+details</div> <div>org.apache.spark.sql.DataFrameReader.json(DataFrameReader.scala:244) sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method) sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:57) sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43) java.lang.reflect.Method.invoke(Method.java:606)</div>	2016/02/14 18:24:03	2 s	2/2	607.4 KB

Display the first 10 records

In [5]:

```
txns.show( 10 )
```

```
+-----+-----+-----+-----+-----+
-----+-----+-----+-----+
|CashOrCredit|      creditCardNo|customerNo|      lineItems| merc
hantCity|      state|      tDate|      txnNo|
+-----+-----+-----+-----+-----+
-----+-----+-----+-----+
|      credit|4971-xxxx-xxxx-5769|      4004819|[[015.82,Team Spo...| Bro
wnsville|      Texas|06-27-2011|00000000|
|      credit|3787-xxxx-xxxx-6017|      4003459|[[089.28,Water Sp...|
Houston|      Texas|02-07-2011|00000001|
|      credit|5951-xxxx-xxxx-4036|      4009112|[[067.51,Exercise...|
Eugene|      Oregon|03-02-2011|00000002|
|      credit|3793-xxxx-xxxx-3180|      4009376|[[043.38,Water Sp...|
Paterson|New Jersey|01-23-2011|00000003|
|      credit|3913-xxxx-xxxx-4556|      4006758|[[193.65,Outdoor ...|
Gresham|      Oregon|05-07-2011|00000004|
|      credit|4629-xxxx-xxxx-3692|      4000951|[[104.47,Exercise...| De
s Moines|      Iowa|12-07-2011|00000005|
|      credit|4032-xxxx-xxxx-1996|      4002494|[[093.97,Jumping,...| St.
Louis | Missouri|05-02-2011|00000006|
|      credit|3551-xxxx-xxxx-0696|      4000599|[[197.33,Exercise...|
Phoenix|      Arizona|06-02-2011|00000007|
|      credit|3282-xxxx-xxxx-5190|      4007057|[[128.98,Winter S...|Overl
and Park|      Kansas|03-06-2011|00000008|
|      credit|4621-xxxx-xxxx-9258|      4005366|[[074.57,Water Sp...|
Fremont|California|06-22-2011|00000009|
+-----+-----+-----+-----+-----+
-----+-----+-----+-----+
only showing top 10 rows
```

In [6]:

```
txns.persist()
```

Out[6]:

```
DataFrame[CashOrCredit: string, creditCardNo: string, customerNo: strin
g, lineItems: array<struct<amount:string,category:string,product:string
>>, merchantCity: string, state: string, tDate: string, txnNo: string]
```

Check the storage. The RDD should have been created.

- Check the size of the RDD. Number of partitions and percentage of cache.

Storage

RDDs

RDD Name	Storage Level	Cached Partitions	Fraction Cached	Size in Memory
Scan JSONRelation[CashOrCredit#8,creditCardNo#9,customerNo#10,lineItems#11,merchantCity#12,state#13,tDate#14,txnNo#15] InputPaths: hdfs://sparklab.awesomestats.in/sparklab/txnjsonsmall	Memory Serialized 1x Replicated	1	50%	84.1 KB

The line items are nested structure in each transaction. Display the lineitems of first 5 transactions

In [7]:

```
txns.select( "lineItems" ).take( 5 )
```

Out[7]:

```
[Row(lineItems=[Row(amount='015.82', category='Team Sports', product='C
heerleading'), Row(amount='086.47', category='Water Sports', product='W
hitewater Rafting'), Row(amount='063.08', category='Exercise & Fitnes
s', product='Gym Mats'), Row(amount='068.80', category='Exercise & Fitn
ess', product='Weightlifting Machine Accessories'), Row(amount='092.4
9', category='Team Sports', product='Lacrosse'), Row(amount='083.92', c
ategory='Outdoor Recreation', product='Lawn Games')]),
Row(lineItems=[Row(amount='089.28', category='Water Sports', product
='Water Tubing'), Row(amount='042.38', category='Water Sports', product
='Surfing'), Row(amount='062.80', category='Team Sports', product='Chee
rleading')]),
Row(lineItems=[Row(amount='067.51', category='Exercise & Fitness', pro
duct='Exercise Bands'), Row(amount='154.57', category='Team Sports', pr
oduct='Rugby'), Row(amount='100.18', category='Outdoor Recreation', pro
duct='Skateboarding'), Row(amount='190.52', category='Exercise & Fitnes
s', product='Foam Rollers'), Row(amount='054.35', category='Water Sport
s', product='Kitesurfing')]),
Row(lineItems=[Row(amount='043.38', category='Water Sports', product
='Boating'), Row(amount='106.27', category='Team Sports', product='Rugb
y'), Row(amount='164.86', category='Combat Sports', product='Fencing'),
Row(amount='164.94', category='Racquet Sports', product='Tennis'), Row
(amount='007.36', category='Exercise & Fitness', product='Gym Mats'), R
ow(amount='110.56', category='Outdoor Recreation', product='Skateboardi
ng')]),
Row(lineItems=[Row(amount='193.65', category='Outdoor Recreation', pro
duct='Deck Shuffleboard'), Row(amount='135.98', category='Winter Sport
s', product='Snowshoeing'), Row(amount='063.27', category='Racquet Spor
ts', product='Racquetball'), Row(amount='151.53', category='Dancing', p
roduct='Ballet Bars'), Row(amount='088.69', category='Gymnastics', prod
uct='Balance Beams'), Row(amount='070.75', category='Outdoor Play Equip
ment', product='Swing Sets')])]
```

Check the next stage for select statement

Stages for All Jobs

Completed Stages: 5

Completed Stages (5)

Stage Id	Description		Submitted	Duration	Tasks: Succeeded/Total	Input
4	take at <ipython-input-13-d303f260d6b0>:1	+details	2016/02/14 18:37:25	41 ms	<div>1/1</div>	84.1 KB
3	showString at NativeMethodAccessorImpl.java:-2	+details	2016/02/14 18:35:22	34 ms	<div>1/1</div>	84.1 KB
2	showString at NativeMethodAccessorImpl.java:-2	+details	2016/02/14 18:34:09	0.8 s	<div>1/1</div>	320.0 KB
1	json at NativeMethodAccessorImpl.java:-2	+details	2016/02/14 18:30:54	0.3 s	<div>2/2</div>	607.4 KB
0	json at NativeMethodAccessorImpl.java:-2	+details	2016/02/14 18:24:03	2 s	<div>2/2</div>	607.4 KB

Also the SQL section for detailed logical plan

SQL

Completed Queries

ID	Description		Submitted	Duration	Jobs	Detail
2	takeAndServe at NativeMethodAccessorImpl.java:-2	+details	2016/02/14 18:37:25	72 ms	4	<div> <div>== Parsed Logical Plan ==</div> <div> -- Parsed Logical Plan -- 'Project [unresolvedalias('lineItems)] +- Relation[CashOrCredit#8,creditCardNo#9,customerNo#10,lineItems#11,merchantCity#12,state#13,tDate#14,txnNo#15] JSONRelation </div> </div> <div> <div>== Analyzed Logical Plan ==</div> <div> lineItems: array<struct<amount:string,category:string,product:string>> Project [lineItems#11] +- Relation[CashOrCredit#8,creditCardNo#9,customerNo#10,lineItems#11,merchantCity#12,state#13,tDate#14,txnNo#15] JSONRelation </div> </div> <div> <div>== Optimized Logical Plan ==</div> <div> Project [lineItems#11] +- InMemoryRelation [CashOrCredit#8,creditCardNo#9,customerNo#10,lineItems#11,merchantCity#12,state#13,tDate#14,txnNo#15], true, 10000, StorageLevel(false, true, false, false, 1), Scan JSONRelation[CashOrCredit#8,creditCardNo#9,customerNo#10,lineItems#11,merchantCity#12,sta </div> </div>

Import the explode function to flatten the records

In [8]:

```
from pyspark.sql.functions import *
```

In [9]:

```
## Explode and flatten the nested structure into a set of columns
txns_new = txns.select( 'txnNo', 'tDate', 'customerNo', 'merchantCity', 'state', 'ite
                        'item.product', 'item.amount',
                        explode( txns.lineItems ).alias( 'item' ) ).drop( 'item')
```

In [10]:

```
# Show 10 records
txns_new.show( 10 )
```

```
+-----+-----+-----+-----+-----+-----+
| txnNo|   tDate|customerNo|merchantCity| state|           category|
product|amount|
+-----+-----+-----+-----+-----+-----+
|00000000|06-27-2011|  4004819| Brownsville| Texas|           Team Sports|
Cheerleading|015.82|
|00000000|06-27-2011|  4004819| Brownsville| Texas|           Water Sports|
Whitewater Rafting|086.47|
|00000000|06-27-2011|  4004819| Brownsville| Texas|Exercise & Fitness|
Gym Mats|063.08|
|00000000|06-27-2011|  4004819| Brownsville| Texas|Exercise & Fitness|
Weightlifting Mac...|068.80|
|00000000|06-27-2011|  4004819| Brownsville| Texas|           Team Sports|
Lacrosse|092.49|
|00000000|06-27-2011|  4004819| Brownsville| Texas|Outdoor Recreation|
Lawn Games|083.92|
|00000001|02-07-2011|  4003459|      Houston| Texas|           Water Sports|
Water Tubing|089.28|
|00000001|02-07-2011|  4003459|      Houston| Texas|           Water Sports|
Surfing|042.38|
|00000001|02-07-2011|  4003459|      Houston| Texas|           Team Sports|
Cheerleading|062.80|
|00000002|03-02-2011|  4009112|      Eugene|Oregon|Exercise & Fitness|
Exercise Bands|067.51|
+-----+-----+-----+-----+-----+-----+
only showing top 10 rows
```

Register the new table as temporary (in memory) table, so that we can run SQL Queries on it

In [11]:

```
# Register the dataframe as an temporary sql table into memory.. so that we can go an
txns_new.registerTempTable("txnrecords")
```

Find revenue generated by state and product

In [12]:

```
revenue_by_state = sqlContext.sql( "select state, product, sum( amount ) as total fro
```

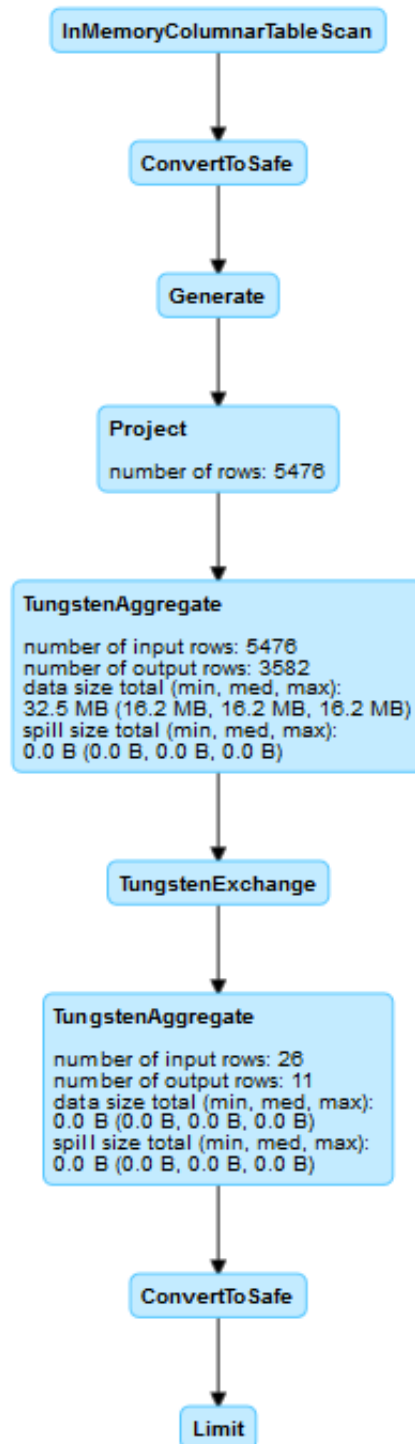
In [13]:

```
# show the first 10 records
revenue_by_state.show( 10 )
```

state	product	total
Oregon	Rugby	261.4
Texas	Parachutes	706.58
Oregon	Scuba Diving & Sn...	264.55
Utah	Wrestling	207.79000000000002
Kentucky	Bobsledding	232.84999999999997
Florida	Foam Rollers	387.27
Massachusetts	Air Hockey	120.66
Alabama	Windsurfing	296.17999999999995
Arizona	Jumping Stilts	96.00999999999999
Pennsylvania	Disc Golf	28.08

only showing top 10 rows

DAG Visualization under SQL



In [14]:

```
# We can also register the result sets as temporary tables
revenue_by_state.registerTempTable('state_revenue')
```

Write an UDF (User defined function) to extract week day name from the date field


```
# Define a user defined function to be invoked from sql query. For example, deriving
import datetime
def getDayOfWeek( date):
    return datetime.datetime.strptime(date, "%m-%d-%Y").strftime('%A')
```

In [16]:

```
# Register the function
from pyspark.sql.types import StringType
sqlContext.udf.register("getDayOfWeek", lambda date: getDayOfWeek( date ), StringType)
```

In [17]:

[illegible]

In [18]:

```
revenue_by_state.show( 10 )
```

weekday	total
Thursday	94549.2
Wednesday	85091.56
Monday	81712.77
Sunday	79634.08
Tuesday	79594.51
Saturday	78114.84
Friday	71809.1

In [19]:

[illegible]

In [20]:

```
cust_df.show( 10 )
```

CustID	FirstName	LastName	Age	Profession
4000001	Kristina	Chung	55	Pilot
4000002	Paige	Chen	74	Teacher
4000003	Sherri	Melton	34	Firefighter
4000004	Gretchen	Hill	66	Computer hardware...
4000005	Karen	Puckett	74	Lawyer
4000006	Patrick	Song	42	Veterinarian
4000007	Elsie	Hamilton	43	Pilot
4000008	Hazel	Bender	63	Carpenter
4000009	Malcolm	Wagner	39	Artist
4000010	Dolores	McLaughlin	60	Writer

only showing top 10 rows

Finding total money spent by each customers

In [21]:

```
top_10_custs = sqlContext.sql( '''select customerNo, round( sum( amount ), 2 )
                                as total from txnrecords group by customerNo
                                order by total desc limit 10''' )
```

In [22]:

```
top_10_custs.registerTempTable( "top_10_custs" )
cust_df.registerTempTable( "custs" )
```

Joining with customer table to find top 10 customers based on total money spent

In [23]:

```
top_10_cust_names = sqlContext.sql( '''select a.CustID, a.FirstName, a.LastName,
                                a.Age, b.total from custs a join top_10_custs b
                                on a.CustID = b.customerNo order by b.total desc''' )
```

```
top_10_cust_names.show( 10 )
```

Timeline



For Executors logs, click on the attempt id on down left

Kill Application						
User:	hadoop					
Name:	pyspark-shell					
Application Type:	SPARK					
Application Tags:						
YarnApplicationState:	RUNNING: AM has registered with RM and started running.					
FinalStatus Reported by AM:	Application has not completed yet.					
Started:	Sun Feb 14 19:43:19 +0100 2016					
Elapsed:	27mins, 25sec					
Tracking URL:	ApplicationMaster					
Diagnostics:						
Total Resource Preempted:	<memory-0, vCores-0>					
Total Number of Non-AM Containers Preempted:	0					
Total Number of AM Containers Preempted:	0					
Resource Preempted from Current Attempt:	<memory-0, vCores-0>					
Number of Non-AM Containers Preempted from Current Attempt:	0					
Aggregate Resource Allocation:	6553939 MB-seconds, 4916 vcore-seconds					

For each Executor logs, click on logs link on each containers

Application Attempt State: RUNNING

AM Container: container_1455474893226_0002_01_000001

Node: 192.168.133.129/0

Tracking URL: ApplicationMaster

Diagnostics Info:

Application Attempt Headroom : <memory:2000, vCores:1>

Application

Total Allocated Containers: 4

Each table cell represents the number of NodeLocal/RackLocal/OffSwitch containers satisfied by NodeLocal/RackLocal/OffSwitch resource requests.

	Node Local Request	Rack Local Request	Off Switch Request
Num Node Local Containers (satisfied by)	0		
Num Rack Local Containers (satisfied by)	0	0	
Num Off Switch Containers (satisfied by)	0		4

Total Outstanding Resource Requests: <memory:0, vCores:0>

Priority	ResourceName	Capability	NumContainers	RelaxLocality	NodeLabelExpression
Show 20	entries				Search:
Container ID	Node		Container Exit Status		Logs
container_1455474893226_0002_01_000003	http://sparklab.awesomestats.in:8042	0			Logs
container_1455474893226_0002_01_000002	http://sparklab.awesomestats.in:8042	0			Logs
container_1455474893226_0002_01_000001	http://sparklab.awesomestats.in:8042	0			Logs

Showing 1 to 3 of 3 entries

First Previous

Make a note of things you learnt in the exercise.