

Spark Developer Training - 3 Days

Manaranjan Pradhan

manaranjan@enablecloud.com

This notebook is given as part of Spark Training to Participants. Forwarding others is strictly prohibited.

Analyze the captains performance in ODIs and Test matches in Cricket

This notebook gives an overview of how to read data into Spark Framework and apply some basic RDD operations

In [1]:

```
# Spark Context Created.. Let's verify it  
sc
```

Out[1]:

```
<pyspark.context.SparkContext at 0x7fa5801ad278>
```

ODI Performance Analysis

Read the capatings_ODI.csv file from local file system

In [2]:

```
captains_odis = sc.textFile( "file:///home/hadoop/lab/data/captains_ODI.csv" )
```

Check the first record

In [3]:

```
captains_odis.first()
```

Out[3]:

```
'Ponting  R T,Australia,1995-2012,230,165,51,14,124'
```

Display the first 10 records

In [4]:

```
captains_odis.take( 10 )
```

Out[4]:

```
['Ponting R T,Australia,1995-2012,230,165,51,14,124',  
'Fleming S P,New Zealand,1994-2007,218,98,106,14,105',  
'Ranatunga A,Sri Lanka,1982-1999,193,89,95,9,102',  
'Dhoni M S*,India,2004-,186,103,68,15,88',  
'Border A R,Australia,1979-1994,178,107,67,4,86',  
'Azharuddin M,India,1985-2000,174,89,77,8,96',  
'Smith G C,South Africa,2002-2013,149,91,51,7,74',  
'Ganguly S C,India,1992-2007,147,76,66,5,74',  
'Cronje W J,South Africa,1992-2000,140,99,37,4,74',  
'Imran Khan,Pakistan,1974-1992,139,75,59,5,70']
```

Checking the RDD type

In [5]:

```
# captains_odis should be of RDD type  
type( captains_odis )
```

Out[5]:

```
pyspark.rdd.RDD
```

Counting the number of records

In [6]:

```
# Count the number of records  
captains_odis.count()
```

Out[6]:

```
98
```

Parsing the lines and creating records

The RDD does not have a header. Let's define a header and then assign it to the RDD

In [7]:

```
fields = ("name", "country", "career", "matches", "won", "lost", "ties", "toss" )
```

In [8]:

```
from collections import namedtuple
```

Lets define a tuple (a record) for each line. We need to iterate through each line and convert that into a record. The record can be defined as a namedtuple type and called Captain. Let's also link the names for the fields

In [9]:

```
Captain = namedtuple( 'Captain', fields )
```

Create a function to parse the lines and create namedtuples

In [10]:

```
# Define a function to parse each line and convert them into records
def parseRecs( line ):
    fields = line.split(",")
    return Captain( fields[0], fields[1], fields[2], int( fields[3] ),
                    int( fields[4] ), int(fields[5]), int(fields[6]), int(fields[7] ) )
```

Iterate throught the data and convert them into records

In [11]:

```
captains = captains_odis.map( lambda rec: parseRecs( rec) )
```

In [12]:

```
# Now captains refer to all the records. Let's display the first 10 records.
captains.take( 10 )
```

Out[12]:

```
[Captain(name='Ponting R T', country='Australia', career='1995-2012',
matches=230, won=165, lost=51, ties=14, toss=124),
 Captain(name='Fleming S P', country='New Zealand', career='1994-2007',
matches=218, won=98, lost=106, ties=14, toss=105),
 Captain(name='Ranatunga A', country='Sri Lanka', career='1982-1999',
matches=193, won=89, lost=95, ties=9, toss=102),
 Captain(name='Dhoni M S*', country='India', career='2004-', matches=186,
won=103, lost=68, ties=15, toss=88),
 Captain(name='Border A R', country='Australia', career='1979-1994', matches=178,
won=107, lost=67, ties=4, toss=86),
 Captain(name='Azharuddin M', country='India', career='1985-2000', matches=174,
won=89, lost=77, ties=8, toss=96),
 Captain(name='Smith G C', country='South Africa', career='2002-2013',
matches=149, won=91, lost=51, ties=7, toss=74),
 Captain(name='Ganguly S C', country='India', career='1992-2007', matches=147,
won=76, lost=66, ties=5, toss=74),
 Captain(name='Cronje W J', country='South Africa', career='1992-2000',
matches=140, won=99, lost=37, ties=4, toss=74),
 Captain(name='Imran Khan', country='Pakistan', career='1974-1992', matches=139,
won=75, lost=59, ties=5, toss=70)]
```

In [13]:

```
# What is the type of the captains RDD
type( captains )
```

Out[13]:

```
pyspark.rdd.PipelinedRDD
```

Filtering records

Filter only those captains that have captained for at least 100 ODI matches. And then we can compare the statistics of these captains

In [14]:

```
captains_100 = captains.filter( lambda rec: rec.matches > 100 )
```

How many captains have captained their country for more than 100 ODIs?

In [15]:

```
captains_100.count()
```

Out[15]:

```
16
```

In [16]:

```
# Who are these captains
captains.take( 10 )
```

Out[16]:

```
[Captain(name='Ponting R T', country='Australia', career='1995-2012',
matches=230, won=165, lost=51, ties=14, toss=124),
 Captain(name='Fleming S P', country='New Zealand', career='1994-2007',
matches=218, won=98, lost=106, ties=14, toss=105),
 Captain(name='Ranatunga A', country='Sri Lanka', career='1982-1999',
matches=193, won=89, lost=95, ties=9, toss=102),
 Captain(name='Dhoni M S*', country='India', career='2004-', matches=186,
won=103, lost=68, ties=15, toss=88),
 Captain(name='Border A R', country='Australia', career='1979-1994', matches=178,
won=107, lost=67, ties=4, toss=86),
 Captain(name='Azharuddin M', country='India', career='1985-2000', matches=174,
won=89, lost=77, ties=8, toss=96),
 Captain(name='Smith G C', country='South Africa', career='2002-2013', matches=149,
won=91, lost=51, ties=7, toss=74),
 Captain(name='Ganguly S C', country='India', career='1992-2007', matches=147,
won=76, lost=66, ties=5, toss=74),
 Captain(name='Cronje W J', country='South Africa', career='1992-2000', matches=140,
won=99, lost=37, ties=4, toss=74),
 Captain(name='Imran Khan', country='Pakistan', career='1974-1992', matches=139,
won=75, lost=59, ties=5, toss=70)]
```

In [17]:

```
# Filtering: Captains who have more wins than losses...
captains_more_wins = captains_100.filter( lambda rec: rec.won > rec.lost )
```

Collecting the results in the driver

In [18]:

```
# Captains with more wins than losses
captains_more_wins.map( lambda rec: rec.name ).collect()
```

Out[18]:

```
['Ponting   R T',
 'Dhoni     M S*',
 'Border    A R',
 'Azharuddin M',
 'Smith     G C',
 'Ganguly   S C',
 'Cronje    W J',
 'Imran Khan',
 'Jayawardene D P M',
 'Jayasuriya S T',
 'Wasim Akram',
 'Waugh     S R',
 'Richards  I V A']
```

In [19]:

```
# Captains with less wins than losses
captains_more_loses = captains_100.filter( lambda rec: rec.won <= rec.lost )
captains_more_loses.map( lambda rec: rec.name ).collect()
```

Out[19]:

```
['Fleming   S P', 'Ranatunga  A', 'Lara      B C']
```

Creating a subset of data by filtering columns

In [20]:

```
# Which country has played how many matches..
countries = captains.map( lambda rec: ( rec.country , rec.matches) )
```

In [21]:

```
countries.take( 10 )
```

Out[21]:

```
[('Australia', 230),  
 ('New Zealand', 218),  
 ('Sri Lanka', 193),  
 ('India', 186),  
 ('Australia', 178),  
 ('India', 174),  
 ('South Africa', 149),  
 ('India', 147),  
 ('South Africa', 140),  
 ('Pakistan', 139)]
```

Aggregating values by keys

In [22]:

```
# Aggregate by countries  
matches_countries = countries.reduceByKey( lambda a, b: a + b )
```

In [23]:

```
matches_countries.take( 20 )
```

Out[23]:

```
[('Pakistan', 781),  
 ('South Africa', 463),  
 ('Ireland', 93),  
 ('Australia', 832),  
 ('West Indies', 658),  
 ('Kenya', 114),  
 ('India', 770),  
 ('Bermuda', 31),  
 ('Netherlands', 31),  
 ('Afghanistan', 50),  
 ('England', 554),  
 ('Canada', 27),  
 ('Zimbabwe', 394),  
 ('Sri Lanka', 710),  
 ('Bangladesh', 251),  
 ('New Zealand', 608)]
```

In [24]:

```
# Sort the countries by the number of matches they played.  
# Sort by names...(sort by key)  
matches_countries.sortByKey().collect()
```

Out[24]:

```
[('Afghanistan', 50),  
 ('Australia', 832),  
 ('Bangladesh', 251),  
 ('Bermuda', 31),  
 ('Canada', 27),  
 ('England', 554),  
 ('India', 770),  
 ('Ireland', 93),  
 ('Kenya', 114),  
 ('Netherlands', 31),  
 ('New Zealand', 608),  
 ('Pakistan', 781),  
 ('South Africa', 463),  
 ('Sri Lanka', 710),  
 ('West Indies', 658),  
 ('Zimbabwe', 394)]
```

Sorting records

In [25]:

```
# Sort the countries by the number of matches they played by descending order..  
# by names (key)  
matches_countries.sortByKey( ascending = False ).collect()
```

Out[25]:

```
[('Zimbabwe', 394),  
 ('West Indies', 658),  
 ('Sri Lanka', 710),  
 ('South Africa', 463),  
 ('Pakistan', 781),  
 ('New Zealand', 608),  
 ('Netherlands', 31),  
 ('Kenya', 114),  
 ('Ireland', 93),  
 ('India', 770),  
 ('England', 554),  
 ('Canada', 27),  
 ('Bermuda', 31),  
 ('Bangladesh', 251),  
 ('Australia', 832),  
 ('Afghanistan', 50)]
```

In [26]:

```
# Sort by values.. default is ascending...
matches_countries.sortBy( lambda rec: rec[1] ).collect()
```

Out[26]:

```
[('Canada', 27),
 ('Bermuda', 31),
 ('Netherlands', 31),
 ('Afghanistan', 50),
 ('Ireland', 93),
 ('Kenya', 114),
 ('Bangladesh', 251),
 ('Zimbabwe', 394),
 ('South Africa', 463),
 ('England', 554),
 ('New Zealand', 608),
 ('West Indies', 658),
 ('Sri Lanka', 710),
 ('India', 770),
 ('Pakistan', 781),
 ('Australia', 832)]
```

In [27]:

```
# Sort by value by descending
matches_countries.sortBy( lambda rec: rec[1], ascending = False ).collect()
```

Out[27]:

```
[('Australia', 832),
 ('Pakistan', 781),
 ('India', 770),
 ('Sri Lanka', 710),
 ('West Indies', 658),
 ('New Zealand', 608),
 ('England', 554),
 ('South Africa', 463),
 ('Zimbabwe', 394),
 ('Bangladesh', 251),
 ('Kenya', 114),
 ('Ireland', 93),
 ('Afghanistan', 50),
 ('Bermuda', 31),
 ('Netherlands', 31),
 ('Canada', 27)]
```


In [28]:

```
# Captains by percentage of wins
captains_100_percent_wins = captains_100.map(
    lambda rec: ( rec.name, round( rec.won/rec.matches, 2 ) ) )

# Sort by percentage wins
captains_100_percent_wins.sortBy(
    lambda rec: rec[1], ascending = False ).collect()
```

Out[28]:

```
[('Ponting R T', 0.72),
 ('Cronje W J', 0.71),
 ('Richards I V A', 0.64),
 ('Waugh S R', 0.63),
 ('Smith G C', 0.61),
 ('Wasim Akram', 0.61),
 ('Border A R', 0.6),
 ('Jayasuriya S T', 0.56),
 ('Dhoni M S*', 0.55),
 ('Jayawardene D P M', 0.55),
 ('Imran Khan', 0.54),
 ('Ganguly S C', 0.52),
 ('Azharuddin M', 0.51),
 ('Lara B C', 0.47),
 ('Ranatunga A', 0.46),
 ('Fleming S P', 0.45)]
```

Exercise for the participants

- Filter countries which have played more than hundred matches
- Sort counties by the percentage of matches they
- Find the top10 lucky captains in terms of TOSS w

In [29]:

```
# Read the Captains_Test.csv file
captains_tests = sc.textFile( "file:///home/hadoop/lab/data/captains_Test.csv" )
```

In [30]:

```
# Parse the records
captains_tests_recs = captains_tests.map( lambda rec: parseRecs( rec ) )
```

```
# Display the first 10 records
captains_tests_recs.take( 10 )
```

```
[Captain(name='Smith G C', country='South Africa', career='2002-2014', matches=109, won=53, lost=29, ties=27, toss=58),
 Captain(name='Border A R', country='Australia', career='1978-1994', matches=93, won=32, lost=22, ties=38, toss=47),
 Captain(name='Fleming S P', country='New Zealand', career='1994-2008', matches=80, won=28, lost=27, ties=25, toss=38),
 Captain(name='Ponting R T', country='Australia', career='1995-2012', matches=77, won=48, lost=16, ties=13, toss=37),
 Captain(name='Lloyd C H', country='West Indies', career='1966-1984', matches=74, won=36, lost=12, ties=26, toss=35),
 Captain(name='Dhoni M S*', country='India', career='2005-', matches=60, won=27, lost=18, ties=15, toss=27),
 Captain(name='Waugh S R', country='Australia', career='1985-2004', matches=57, won=41, lost=9, ties=7, toss=31),
 Captain(name='Ranatunga A', country='Sri Lanka', career='1982-2000', matches=56, won=12, lost=19, ties=25, toss=29),
 Captain(name='Atherton M A', country='England', career='1989-2001', matches=54, won=13, lost=21, ties=20, toss=23),
 Captain(name='Cronje W J', country='South Africa', career='1992-2000', matches=53, won=27, lost=11, ties=15, toss=22)]
```

[illegible]

```
## How many captains?
captains tests 100.take( 10 )
```

```
[Captain(name='Smith G C', country='South Africa', career='2002-2014', matches=109, won=53, lost=29, ties=27, toss=58)]
```

[illegible]

In [35]:

```
captains_tests_50.take( 10 )
```

Out[35]:

```
[Captain(name='Smith G C', country='South Africa', career='2002-2014', matches=109, won=53, lost=29, ties=27, toss=58),
 Captain(name='Border A R', country='Australia', career='1978-1994', matches=93, won=32, lost=22, ties=38, toss=47),
 Captain(name='Fleming S P', country='New Zealand', career='1994-2008', matches=80, won=28, lost=27, ties=25, toss=38),
 Captain(name='Ponting R T', country='Australia', career='1995-2012', matches=77, won=48, lost=16, ties=13, toss=37),
 Captain(name='Lloyd C H', country='West Indies', career='1966-1984', matches=74, won=36, lost=12, ties=26, toss=35),
 Captain(name='Dhoni M S*', country='India', career='2005-', matches=60, won=27, lost=18, ties=15, toss=27),
 Captain(name='Waugh S R', country='Australia', career='1985-2004', matches=57, won=41, lost=9, ties=7, toss=31),
 Captain(name='Ranatunga A', country='Sri Lanka', career='1982-2000', matches=56, won=12, lost=19, ties=25, toss=29),
 Captain(name='Atherton M A', country='England', career='1989-2001', matches=54, won=13, lost=21, ties=20, toss=23),
 Captain(name='Cronje W J', country='South Africa', career='1992-2000', matches=53, won=27, lost=11, ties=15, toss=22)]
```

In [36]:

```
# Sort the captains by percentage of wins
captain_top = captains_tests_50.map(
    lambda rec: ( rec.name,
                  round( rec.won/rec.matches,
                        2 ) ) ).sortBy( lambda rec: rec[1], ascending = False )
```

In [37]:

```
captain_top.collect()
```

Out[37]:

```
[('Waugh S R', 0.72),
 ('Ponting R T', 0.62),
 ('Cronje W J', 0.51),
 ('Vaughan M P', 0.51),
 ('Smith G C', 0.49),
 ('Lloyd C H', 0.49),
 ('Dhoni M S*', 0.45),
 ('Fleming S P', 0.35),
 ('Border A R', 0.34),
 ('Atherton M A', 0.24),
 ('Ranatunga A', 0.21)]
```

Joining multiple data sets

In [38]:

```
# Lets join both ODI and Test captaincy details.  
# Default is inner join...  
all_time_best_captains = captains_100_percent_wins.join( captain_top )
```

In [39]:

```
all_time_best_captains.collect()
```

Out[39]:

```
[('Smith G C', (0.61, 0.49)),  
 ('Border A R', (0.6, 0.34)),  
 ('Fleming S P', (0.45, 0.35)),  
 ('Cronje W J', (0.71, 0.51)),  
 ('Ponting R T', (0.72, 0.62)),  
 ('Ranatunga A', (0.46, 0.21)),  
 ('Dhoni M S*', (0.55, 0.45)),  
 ('Waugh S R', (0.63, 0.72))]
```

In [40]:

```
## Best by test match wins  
all_time_best_captains.sortBy( lambda rec: rec[1][1],  
                               ascending = False ).collect()
```

Out[40]:

```
[('Waugh S R', (0.63, 0.72)),  
 ('Ponting R T', (0.72, 0.62)),  
 ('Cronje W J', (0.71, 0.51)),  
 ('Smith G C', (0.61, 0.49)),  
 ('Dhoni M S*', (0.55, 0.45)),  
 ('Fleming S P', (0.45, 0.35)),  
 ('Border A R', (0.6, 0.34)),  
 ('Ranatunga A', (0.46, 0.21))]
```

In [41]:

```
## Best by ODI match wins  
all_time_best_captains.sortBy( lambda rec: rec[1][0],  
                               ascending = False ).collect()
```

Out[41]:

```
[('Ponting R T', (0.72, 0.62)),  
 ('Cronje W J', (0.71, 0.51)),  
 ('Waugh S R', (0.63, 0.72)),  
 ('Smith G C', (0.61, 0.49)),  
 ('Border A R', (0.6, 0.34)),  
 ('Dhoni M S*', (0.55, 0.45)),  
 ('Ranatunga A', (0.46, 0.21)),  
 ('Fleming S P', (0.45, 0.35))]
```

In [42]:

```
## Now let's flatten the structure and store the results into a file...
best_captains = all_time_best_captains.map( lambda rec:
                                            ( rec[0],
                                              rec[1][0],
                                              rec[1][1] ) )

best_captains.take( 10 )
```

Out[42]:

```
[('Smith   G C', 0.61, 0.49),
 ('Border  A R', 0.6, 0.34),
 ('Fleming  S P', 0.45, 0.35),
 ('Cronje   W J', 0.71, 0.51),
 ('Ponting  R T', 0.72, 0.62),
 ('Ranatunga A', 0.46, 0.21),
 ('Dhoni    M S*', 0.55, 0.45),
 ('Waugh    S R', 0.63, 0.72)]
```

Saving results into filesystem

In [43]:

```
best_captains.saveAsTextFile( "file:///home/hadoop/lab/results/captains")
```

Check results

Go to you VM's putty terminal and go to the results directory

```
cd /home/hadoop/lab/results/captains/
```

```
cat part-0000*
```

Saving results into one file. Forcing the bring all results into one partition

In [44]:

```
## Consolidate into one partition
best_captains.repartition( 1 ).
    saveAsTextFile( "file:///home/hadoop/lab/results/bestcaptains")
```

Check results

Go to you VM's putty terminal and go to the results directory

```
cd /home/hadoop/lab/results/bestcaptains/
```

```
cat part-00000
```

Make note of lessons learnt in this lab

