

Spark Developer Training - 3 Days

Manaranjan Pradhan

manaranjan@enablecloud.com

This notebook is given as part of Spark Training to Participants. Forwarding others is strictly prohibited.

Lab: Working with HDFS & Yarn - Retail Data Analysis

Things to learn

- Reading from HDFS
- Reading from MySQL (from RDBMS using JDBC)
- Working with JSON Data - Reading and Parsing
- Working with Spark SQLs
- Applying data transformation using Spark SQL Statements

In [1]:

```
sc
```

Out[1]:

```
<pyspark.context.SparkContext at 0x7f509021d1d0>
```

This Spark Application is running on YARN. So, open the YARN Resource manager UI and verify if the application is running

- To open resource manager web UI, enter <http://hadooplab.bigdataleap.com:8088/>
(<http://hadooplab.bigdataleap.com:8088/>)

Initialize SQLContext from SparkContext

In [2]:

```
from pyspark.sql import SQLContext  
sqlContext = SQLContext(sc)
```

Reading JSON file from HDFS

In [3]:

```
## Read the json file from HDFS
txns = sqlContext.read.json(
    "hdfs://hadooplab.bigdataleap.com/sparklab/txnjsonsmall")
```

Display the first 10 records

In [4]:

```
txns.show( 10 )
```

```
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
|CashOrCredit|creditCardNo|customerNo|lineItems|merc
hantCity|state|tDate|txnNo|
+-----+-----+-----+-----+
+-----+-----+-----+-----+
|credit|4971-xxxx-xxxx-5769|4004819|[[015.82,Team Spo...|Bro
wnsville|Texas|06-27-2011|00000000|
|credit|3787-xxxx-xxxx-6017|4003459|[[089.28,Water Sp...|
Houston|Texas|02-07-2011|00000001|
|credit|5951-xxxx-xxxx-4036|4009112|[[067.51,Exercise...|
Eugene|Oregon|03-02-2011|00000002|
|credit|3793-xxxx-xxxx-3180|4009376|[[043.38,Water Sp...|
Paterson|New Jersey|01-23-2011|00000003|
|credit|3913-xxxx-xxxx-4556|4006758|[[193.65,Outdoor ...|
Gresham|Oregon|05-07-2011|00000004|
|credit|4629-xxxx-xxxx-3692|4000951|[[104.47,Exercise...|De
s Moines|Iowa|12-07-2011|00000005|
|credit|4032-xxxx-xxxx-1996|4002494|[[093.97,Jumping,...|St.
Louis|Missouri|05-02-2011|00000006|
|credit|3551-xxxx-xxxx-0696|4000599|[[197.33,Exercise...|
Phoenix|Arizona|06-02-2011|00000007|
|credit|3282-xxxx-xxxx-5190|4007057|[[128.98,Winter S...|Overl
and Park|Kansas|03-06-2011|00000008|
|credit|4621-xxxx-xxxx-9258|4005366|[[074.57,Water Sp...|
Fremont|California|06-22-2011|00000009|
+-----+-----+-----+-----+
+-----+-----+-----+-----+
only showing top 10 rows
```

The line items are nested structure in each transaction. Display the lineitems of first 5 transactions

In [5]:

```
txns.select( "lineItems" ).take( 5 )
```

Out[5]:

```
[Row(lineItems=[Row(amount='015.82', category='Team Sports', product='Cheerleading'), Row(amount='086.47', category='Water Sports', product='Whitewater Rafting'), Row(amount='063.08', category='Exercise & Fitness', product='Gym Mats'), Row(amount='068.80', category='Exercise & Fitness', product='Weightlifting Machine Accessories'), Row(amount='092.49', category='Team Sports', product='Lacrosse'), Row(amount='083.92', category='Outdoor Recreation', product='Lawn Games')]), Row(lineItems=[Row(amount='089.28', category='Water Sports', product='Water Tubing'), Row(amount='042.38', category='Water Sports', product='Surfing'), Row(amount='062.80', category='Team Sports', product='Cheerleading')]), Row(lineItems=[Row(amount='067.51', category='Exercise & Fitness', product='Exercise Bands'), Row(amount='154.57', category='Team Sports', product='Rugby'), Row(amount='100.18', category='Outdoor Recreation', product='Skateboarding'), Row(amount='190.52', category='Exercise & Fitness', product='Foam Rollers'), Row(amount='054.35', category='Water Sports', product='Kitesurfing')]), Row(lineItems=[Row(amount='043.38', category='Water Sports', product='Boating'), Row(amount='106.27', category='Team Sports', product='Rugby'), Row(amount='164.86', category='Combat Sports', product='Fencing'), Row(amount='164.94', category='Racquet Sports', product='Tennis'), Row(amount='007.36', category='Exercise & Fitness', product='Gym Mats'), Row(amount='110.56', category='Outdoor Recreation', product='Skateboarding')]), Row(lineItems=[Row(amount='193.65', category='Outdoor Recreation', product='Deck Shuffleboard'), Row(amount='135.98', category='Winter Sports', product='Snowshoeing'), Row(amount='063.27', category='Racquet Sports', product='Racquetball'), Row(amount='151.53', category='Dancing', product='Ballet Bars'), Row(amount='088.69', category='Gymnastics', product='Balance Beams'), Row(amount='070.75', category='Outdoor Play Equipment', product='Swing Sets')])])]
```

Import the explode function to flatten the records

In [6]:

```
from pyspark.sql.functions import *
```

In [7]:

```
## Explode and flatten the nested structure into a set of columns
txns_new = txns.select( 'txnNo', 'tDate', 'customerNo', 'merchantCity',
                        'state', 'item.category',
                        'item.product', 'item.amount',
                        explode( txns.lineItems ).alias( 'item' ) ).drop( 'item')
```

In [8]:

```
# Show 10 records
txns_new.show( 10 )
```

```
+-----+-----+-----+-----+-----+-----+
-----+-----+
|  txnNo|    tDate|customerNo|merchantCity| state|          category|
product|amount|
+-----+-----+-----+-----+-----+-----+
-----+-----+
|00000000|06-27-2011|  4004819| Brownsville| Texas|          Team Sports|
Cheerleading|015.82|
|00000000|06-27-2011|  4004819| Brownsville| Texas|          Water Sports|
Whitewater Rafting|086.47|
|00000000|06-27-2011|  4004819| Brownsville| Texas|Exercise & Fitness|
Gym Mats|063.08|
|00000000|06-27-2011|  4004819| Brownsville| Texas|Exercise & Fitness|
Weightlifting Mac...|068.80|
|00000000|06-27-2011|  4004819| Brownsville| Texas|          Team Sports|
Lacrosse|092.49|
|00000000|06-27-2011|  4004819| Brownsville| Texas|Outdoor Recreation|
Lawn Games|083.92|
|00000001|02-07-2011|  4003459|      Houston| Texas|          Water Sports|
Water Tubing|089.28|
|00000001|02-07-2011|  4003459|      Houston| Texas|          Water Sports|
Surfing|042.38|
|00000001|02-07-2011|  4003459|      Houston| Texas|          Team Sports|
Cheerleading|062.80|
|00000002|03-02-2011|  4009112|      Eugene|Oregon|Exercise & Fitness|
Exercise Bands|067.51|
+-----+-----+-----+-----+-----+-----+
-----+-----+
only showing top 10 rows
```

Register the new table as temporary (in memory) table, so that we can run SQL Queries on it

In [9]:

```
# Register the dataframe as an temporary sql table into memory..
# so that we can go and run some sql query
txns_new.registerTempTable("txnrecords")
```

Find revenue generated by state and product

In [10]:

```
revenue_by_state = sqlContext.sql( '''select state, product, sum( amount ) as
                                     total from txnrecords group by state, product''' )
```

In [11]:

```
# show the first 10 records
revenue_by_state.show( 10 )
```

state	product	total
Minnesota	Lawn Games	245.45
Kentucky	Bobsledding	232.84999999999997
Maryland	Playhouses	122.14
Oregon	Camping & Backpac...	88.38
Florida	Foam Rollers	387.27
Oregon	Scuba Diving & Sn...	264.55
Nebraska	Deck Shuffleboard	317.43
Oregon	Rugby	261.4
Utah	Wrestling	207.79000000000002
Texas	Parachutes	706.58

only showing top 10 rows

In [12]:

```
# We can also register the result sets as temporary tables
revenue_by_state.registerTempTable('state_revenue')
```

Write an UDF (User defined function) to extract week day name from the date field

In [13]:

```
# Define a user defined function to be invoked from sql query.
# For example, deriving weekday name from the date field
import datetime
def getDayOfWeek( date):
    return datetime.datetime.strptime(date, "%m-%d-%Y").strftime('%A')
```

Register the function to SQL Context as new UDF

In [14]:

```
# Register the function
from pyspark.sql.types import StringType
sqlContext.udf.register("getDayOfWeek", lambda date: getDayOfWeek( date ),
                        StringType() )
```

Invoke the UDF from the SQL

In [15]:

```
# Write a query to invoke the user defined function..  
# Calculate the total revenue by different weekdays...  
revenue_by_state = sqlContext.sql( '''select weekday as weekday,  
                                     round( sum( amount ), 2 ) as total  
                                     from ( select getDayOfWeek( tDate ) as weekday,  
                                               amount from txnrecords ) txns  
                                     group by weekday order by total desc''' )
```

In [16]:

```
revenue_by_state.show( 10 )
```

```
+-----+-----+  
| weekday|  total|  
+-----+-----+  
| Thursday| 94549.2|  
| Wednesday|85091.56|  
|  Monday|81712.77|  
|  Sunday|79634.08|  
|  Tuesday|79594.51|  
| Saturday|78114.84|  
|  Friday| 71809.1|  
+-----+-----+
```

Reading data from MySql

Check MySql Table

- Go to linux prompt of your VM
- Enter "mysql -u root -p"
- Enter password
- Select Database retail
 - use retail;
- Show tables
 - show tables;
 - describe customers;
- Load data into customers table
LOAD DATA LOCAL INFILE '/home/hadoop/lab/data/custs' INTO TABLE customers FIELDS
TERMINATED BY ',' LINES TERMINATED BY '\r\n';
- List some of the records
 - select * from customers limit 100;

Using jdbc to read from mysql table

In [18]:

```
## Read customer information from mysql table....
cust_df = sqlContext.read.format('jdbc') \
    .options(url='jdbc:mysql://localhost/retail?user=root&password=hadoop123',
             dbtable='customers').load()
```

In [19]:

```
cust_df.show( 10 )
```

```
+-----+-----+-----+---+-----+
| CustID|FirstName|  LastName|Age|          Profession|
+-----+-----+-----+---+-----+
|4000001| Kristina|    Chung| 55|          Pilot|
|4000002|   Paige|    Chen| 74|          Teacher|
|4000003|  Sherri|   Melton| 34|    Firefighter|
|4000004| Gretchen|    Hill| 66|Computer hardware...|
|4000005|   Karen|   Puckett| 74|          Lawyer|
|4000006| Patrick|    Song| 42|    Veterinarian|
|4000007|   Elsie| Hamilton| 43|          Pilot|
|4000008|   Hazel|   Bender| 63|        Carpenter|
|4000009| Malcolm|   Wagner| 39|          Artist|
|4000010| Dolores|McLaughlin| 60|          Writer|
+-----+-----+-----+---+-----+
only showing top 10 rows
```

Finding total money spent by each customers

In [20]:

```
top_10_custs = sqlContext.sql( '''select customerNo, round( sum( amount ), 2 )
                                as total from txnrecords group by customerNo
                                order by total desc limit 10''' )
```

In [21]:

```
top_10_custs.registerTempTable( "top_10_custs" )
cust_df.registerTempTable( "custs_rec" )
```

Joining with customer table to find top 10 customers based on total money spent

In [22]:

```
top_10_cust_names = sqlContext.sql( '''select a.CustID, a.FirstName, a.LastName,
                                           a.Age, b.total from custs_rec a join top_10_custs b
                                           on a.CustID = b.customerNo order by b.total desc''' )
```

In [23]:

```
top_10_cust_names.show( 10 )
```

CustID	FirstName	LastName	Age	total
4007510	Kristin	Levin	73	2204.79
4003293	Martha	Warner	45	2024.67
4003971	Donald	Lamm	34	1869.43
4004260	Courtney	Rubin	54	1869.12
4001058	Gloria	Matthews	53	1791.99
4008914	Samantha	Batchelor	41	1652.06
4004491	Rita	Parks	44	1649.74
4007168	Carolyn	Han	52	1610.76
4001253	Peter	McNamara	74	1516.77
4009672	Samuel	Kidd	61	1485.23

Exercises

- Find out top 5 selling products in each category
- Find top 10 customers for every month

Make a note of things you learnt in the exercise.